

Charlie Snell – csnell22@berkeley.edu

Education

UC Berkeley - *P.h.D. Student, EECS Department [July 2022-May 2026] (Expected)*

Advisors: Dan Klein and Sergey Levine

UC Berkeley (GPA 3.989 out of 4) - *B.A. CS [August 2018-May 2022]*

A+ in: CS288 (natural language processing), CS170 (algorithms) [2nd of 400+ in class final algorithms contest], EE126 (probability), CS188 (artificial intelligence), Physics 137A (quantum mechanics), CS61C (computer architecture), Math 53 (vector calculus), CS61A (CS intro)

Industry Experience

Google Deepmind - *[May 2023-Current]*

Student Researcher

Hosts: Kelvin Xu and Jaehoon Lee

Amazon Web Services - *[Summer 2020]*

Software Engineering Intern

Zelis Healthcare - *[Summer 2019]*

Software Engineering Intern

Publications

The False Promise of Imitating Proprietary LLMs

Arnav Gudibande*, Eric Wallace*, **Charlie Snell***, Xinyang Geng, Hao Liu, Pieter Abbeel, Sergey Levine, Dawn Song
ICLR, 2024.

LMRL Gym: Benchmarks for Multi-Turn Reinforcement Learning with Language Models

Marwa Abdulhai, Isadora White, **Charlie Snell**, Charles Sun, Joey Hong, Yuxiang Zhai, Kelvin Xu, Sergey Levine
arXiv Preprint, 2023.

Learning by Distilling Context

Charlie Snell, Dan Klein, Ruiqi Zhong
arXiv Preprint, 2022.

Offline RL for Natural Language Generation with Implicit Language Q Learning

Charlie Snell, Ilya Kostrikov, Yi Su, Mengjiao Yang, Sergey Levine
ICLR, 2023.

Non-Programmers Can Label Programs Indirectly via Active Examples: A Case Study with Text-to-SQL

Ruiqi Zhong*, **Charlie Snell***, Dan Klein, Jason Eisner
EMNLP, 2023.

Describing Differences between Text Distributions with Natural Language

Ruiqi Zhong, **Charlie Snell**, Dan Klein, Jacob Steinhardt
ICML, 2022.

Context-Aware Language Modeling for Goal-Oriented Dialogue Systems

Charlie Snell, Mengjiao Yang, Justin Fu, Yi Su, Sergey Levine
NAACL 2022, Findings.

Approximating How Single Head Attention Learns

Charlie Snell*, Ruiqi Zhong*, Dan Klein, Jacob Steinhardt
arXiv Preprint, 2021.

The Omniglot Jr. challenge; Can a model achieve child-level character generation and classification?

Eliza Kosoy, Masha Belyi, **Charlie Snell**, Josh Tenenbaum, Brenden Lake, Alison Gopnik
NeurIPS Workshop on BabyMind, 2020.

Invited Talks

G-Research ML Seminar, 2023. *The False Promise of Imitating Proprietary LLMs*

Cohere for AI, 2022. *Learning by Distilling Context*

Carper AI Seminar, 2022. *Offline RL for Natural Language Generation with Implicit Language Q Learning*

Service

Berkeley PhD EECS Admissions - [2023]

Head of Blog at Machine Learning at Berkeley - [January 2021-August 2021]

Teaching for Machine Learning at Berkeley - [January 2020-December 2020]

Reviewer for: ICML 2024, ICLR 2024, COLM 2024, ICLR 2024 Workshop ME-FoMo, ICML 2023, NeurIPS 2023 Workshop Instruction, NeurIPS 2022, NeurIPS 2022 Workshop Offline RL

Awards

UC Berkeley Dean's List - [2019, 2020, 2021]

Honorable Mention for the 2022 CRA outstanding undergraduate researcher award - [2022]